

pg_stat_statements插件分析

函数 `_PG_init` - 负责插件的初始化

函数原型

- 源文件

```
src/backend/utils/fmgr/dfmgr.c
```

- 原型定义

```
/* signatures for PostgreSQL-specific library init/fini functions */
typedef void (*PG_init_t) (void);
typedef void (*PG_fini_t) (void);
```

函数功能

- 插件都必须定义函数 `_PG_init`，PostgreSQL在加载插件对应的动态库文件时，会查找初始化函数 `_PG_init`，该函数负责创建的初始化。以下是PostgreSQL源码中与函数 `_PG_init` 相关的文件。

```
[root@uat-racdb01 postgresql-12.3]# grep -r _PG_init --include='*.c' src/
src/test/modules/test_rls_hooks/test_rls_hooks.c: void _PG_init(void);
src/test/modules/test_rls_hooks/test_rls_hooks.c: _PG_init(void)
src/test/modules/worker_spi/worker_spi.c: void _PG_init(void);
src/test/modules/worker_spi/worker_spi.c: _PG_init(void)
src/test/modules/dummy_seclabel/dummy_seclabel.c: void _PG_init(void);
src/test/modules/dummy_seclabel/dummy_seclabel.c: _PG_init(void)
src/test/modules/test_shm_mq/test.c: void _PG_init(void);
src/backend/storage/lmgr/lwlock.c: * This is only useful for extensions if called from the _PG_init hook
src/backend/storage/ipc/ipci.c: * This is only useful if called from the _PG_init hook of a library that
src/backend/replication/libpqwalreceiver/libpqwalreceiver.c: void _PG_init(void);
src/backend/replication/libpqwalreceiver/libpqwalreceiver.c: _PG_init(void)
src/backend/utils/fmgr/dfmgr.c: * If the library has a _PG_init() function, call it.
src/backend/dfmgr.c: PG_init = (PG_init_t) dlsym(file_scanner->handle, "_PG_init");
src/backend/postmaster/bgworker.c: * This can only be called directly from postmaster or in the _PG_init
src/pl/plpgsql/src/pl_handler.c: * _PG_init() - library load-time initialization
src/pl/plpgsql/src/pl_handler.c: _PG_init(void)
src/pl/perl/plperl.c: void _PG_init(void);
src/pl/perl/plperl.c: _PG_init() - library load-time initialization
src/pl/perl/plperl.c: _PG_init(void)
src/pl/perl/plperl.c: * Does not fully undo the actions of _PG_init() nor make it callable again.
src/pl/tcl/pltcl.c: void _PG_init(void);
src/pl/tcl/pltcl.c: _PG_init() - library load-time initialization
src/pl/tcl/pltcl.c: _PG_init(void)
src/pl/python/plpy_main.c: extern void _PG_init(void);
src/pl/python/plpy_main.c: _PG_init(void)
[root@uat-racdb01 postgresql-12.3]# grep -r \"_PG_init\" --include='*.c' src/
src/backend/dfmgr.c: PG_init = (PG_init_t) dlsym(file_scanner->handle, "_PG_init");
[root@uat-racdb01 postgresql-12.3]#
```

- PostgreSQL在加载插件对应的库文件时，会调用函数 `_PG_init` 完成创建的初始化。

- 函数 `internal_load_library`
- 源文件 `src/backend/utils/fmgr/dfmgr.c`
- 初始化代码

```
/*
 * If the library has a _PG_init() function, call it.
 */
PG_init = (PG_init_t) dlsym(file_scanner->handle, "_PG_init");
if (PG_init)
    (*PG_init) ();
```

- 插件 `pg_stat_statements` 在该函数中通过调用以下几个函数定义一些GUC参数(整型、枚举型和布尔型)，并初始化以下几个pg钩子。
 - 创建GUC变量 (postgresql参数)

```
DefineCustomIntVariable
DefineCustomEnumVariable
DefineCustomBoolVariable
```

- ◦ 安装钩子 `hook`

```
/*
 * Install hooks.
 */
prev_shmem_startup_hook = shmem_startup_hook; /* 保存原来创建共享内存的
hook, 比如其他创建创建的. */
shmem_startup_hook = pgss_shmem_startup; /* 创建创建的共享内存 */
prev_post_parse_analyze_hook = post_parse_analyze_hook;
post_parse_analyze_hook = pgss_post_parse_analyze;
prev_ExecutorStart = ExecutorStart_hook; /* 生成执行计划的hook */
ExecutorStart_hook = pgss_ExecutorStart;
prev_ExecutorRun = ExecutorRun_hook; /* 执行sql语句. */
ExecutorRun_hook = pgss_ExecutorRun;
prev_ExecutorFinish = ExecutorFinish_hook;
ExecutorFinish_hook = pgss_ExecutorFinish;
prev_ExecutorEnd = ExecutorEnd_hook;
ExecutorEnd_hook = pgss_ExecutorEnd;
prev_ProcessUtility = ProcessUtility_hook; /* 非DML语句的执行钩子, 比
如create table、alter table语句. */
ProcessUtility_hook = pgss_ProcessUtility;
```

函数代码

```
/*
 * Module load callback
 */
void
_PG_init(void)
{
    /*
     * In order to create our shared memory area, we have to be loaded via
     * shared_preload_libraries. If not, fall out without hooking into any of
     * the main system. (We don't throw error here because it seems useful to
     * allow the pg_stat_statements functions to be created even when the
     * module isn't active. The functions must protect themselves against
     * being called then, however.)
     */
    if (!process_shared_preload_libraries_in_progress)
        return;

    /*
     * Define (or redefine) custom GUC variables.
     */
    DefineCustomIntVariable("pg_stat_statements.max",
        "Sets the maximum number of statements tracked by
pg_stat_statements.",
        NULL,
        &pgss_max,
        5000,
        100,
        INT_MAX,
        PGC_POSTMASTER,
        0,
        NULL,
        NULL,
        NULL);

    DefineCustomEnumVariable("pg_stat_statements.track",
        "Selects which statements are tracked by pg_stat_statements.",
        NULL,
        &pgss_track,
        PGSS_TRACK_TOP,
        track_options,
        PGC_SUSET,
        0,
        NULL,
```

```

        NULL,
        NULL);

DefineCustomBoolVariable("pg_stat_statements.track_utility",
    "Selects whether utility commands are tracked by
pg_stat_statements.",
    NULL,
    &pgss_track_utility,
    true,
    PGC_SUSET,
    0,
    NULL,
    NULL,
    NULL);

DefineCustomBoolVariable("pg_stat_statements.save",
    "Save pg_stat_statements statistics across server shutdowns.",
    NULL,
    &pgss_save,
    true,
    PGC_SIGHUP,
    0,
    NULL,
    NULL,
    NULL);

EmitWarningsOnPlaceholders("pg_stat_statements");

/*
 * Request additional shared resources. (These are no-ops if we're not in
 * the postmaster process.) We'll allocate or attach to the shared
 * resources in pgss_shmem_startup().
 */
RequestAddinShmemSpace(pgss_memsize());
RequestNamedLWLockTranche("pg_stat_statements", 1);

/*
 * Install hooks.
 */
prev_shmem_startup_hook = shmem_startup_hook;
shmem_startup_hook = pgss_shmem_startup;
prev_post_parse_analyze_hook = post_parse_analyze_hook;
post_parse_analyze_hook = pgss_post_parse_analyze;
prev_ExecutorStart = ExecutorStart_hook;
ExecutorStart_hook = pgss_ExecutorStart;
prev_ExecutorRun = ExecutorRun_hook;
ExecutorRun_hook = pgss_ExecutorRun;
prev_ExecutorFinish = ExecutorFinish_hook;
ExecutorFinish_hook = pgss_ExecutorFinish;

```

```

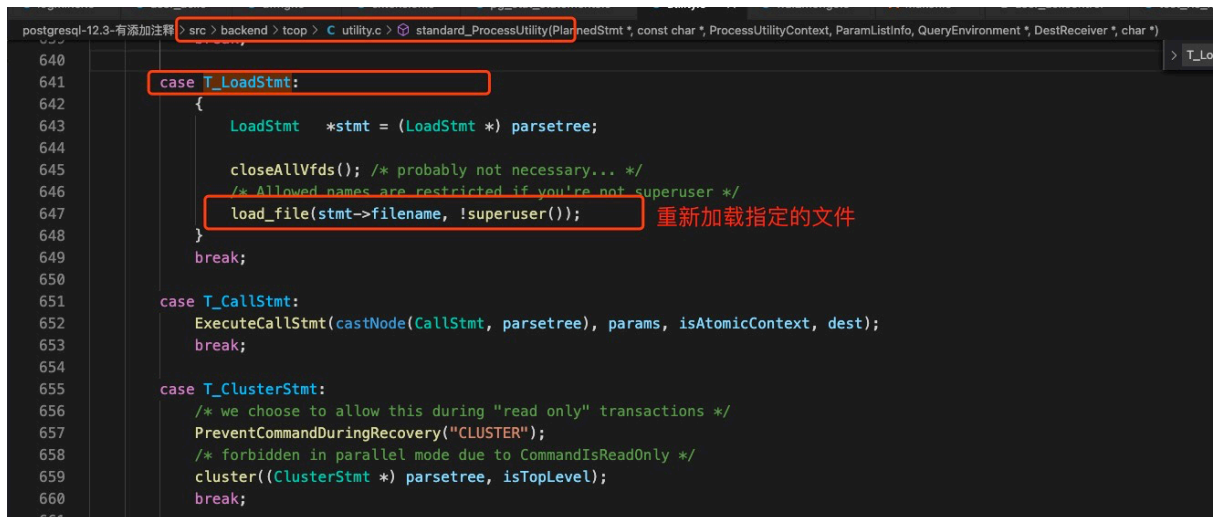
prev_ExecutorEnd = ExecutorEnd_hook;
ExecutorEnd_hook = pgss_ExecutorEnd;
prev_ProcessUtility = ProcessUtility_hook;
ProcessUtility_hook = pgss_ProcessUtility;
}

```

函数 `_PG_fini` - 负责插件的清理

注意事项

- 虽然函数 `_PG_fini` 负责插件的清理工作，但是在使用sql命令 `drop extension` 删除创建时，PG并不会调用该函数。
- 程序执行到以下代码才会调用该函数。但是调试 `pg_ctl reload` 命令，发现并不会执行以下代码。只有sql命令 `load 'user_acl'` 加载库文件时才会触发。



```

640
641 case T_LoadStmt:
642 {
643     LoadStmt *stmt = (LoadStmt *) parsetree;
644
645     closeAllVfds(); /* probably not necessary... */
646     /* Allowed names are restricted if you're not superuser */
647     load_file(stmt->filename, !superuser()); 重新加载指定的文件
648 }
649 break;
650
651 case T_CallStmt:
652     ExecuteCallStmt(castNode(CallStmt, parsetree), params, isAtomicContext, dest);
653     break;
654
655 case T_ClusterStmt:
656     /* we choose to allow this during "read only" transactions */
657     PreventCommandDuringRecovery("CLUSTER");
658     /* forbidden in parallel mode due to CommandIsReadOnly */
659     cluster((ClusterStmt *) parsetree, isTopLevel);
660     break;
661

```

- load命令加载库文件
 - 会话1.

```

postgres=# select pg_backend_pid();
pg_backend_pid
-----
44478
(1 row)

postgres=# load 'user_acl';

```

- gdb调试

```
(gdb) bt
#0  standard_ProcessUtility (pstmt=0x2ec0dd8, queryString=0x2ec0068
"load 'user_acl';", context=PROCESS_UTILITY_TOPLEVEL, params=0x0,
queryEnv=0x0, dest=0x2ec0ed0,
    completionTag=0x7ffde1121260 "") at utility.c:643
#1  0x00007felc3ac8bd0 in uacl_ProcessUtility (pstmt=0x2ec0dd8,
queryString=0x2ec0068 "load 'user_acl';",
context=PROCESS_UTILITY_TOPLEVEL, params=0x0, queryEnv=0x0,
    dest=0x2ec0ed0, completionTag=0x7ffde1121260 "") at user_acl.c:110
#2  0x00007felc38c0c64 in pgss_ProcessUtility (pstmt=0x2ec0dd8,
queryString=0x2ec0068 "load 'user_acl';",
context=PROCESS_UTILITY_TOPLEVEL, params=0x0, queryEnv=0x0,
    dest=0x2ec0ed0, completionTag=0x7ffde1121260 "") at
pg_stat_statements.c:1002
#3  0x00000000008e5866 in ProcessUtility (pstmt=0x2ec0dd8,
queryString=0x2ec0068 "load 'user_acl';",
context=PROCESS_UTILITY_TOPLEVEL, params=0x0, queryEnv=0x0,
    dest=0x2ec0ed0, completionTag=0x7ffde1121260 "") at utility.c:356
#4  0x00000000008e48a2 in PortalRunUtility (portal=0x2f5a238,
pstmt=0x2ec0dd8, isTopLevel=true, setHoldSnapshot=false,
dest=0x2ec0ed0, completionTag=0x7ffde1121260 "")
    at pquery.c:1175
#5  0x00000000008e4aba in PortalRunMulti (portal=0x2f5a238,
isTopLevel=true, setHoldSnapshot=false, dest=0x2ec0ed0,
altdest=0x2ec0ed0, completionTag=0x7ffde1121260 "")
    at pquery.c:1321
#6  0x00000000008e3fff in PortalRun (portal=0x2f5a238,
count=9223372036854775807, isTopLevel=true, run_once=true,
dest=0x2ec0ed0, altdest=0x2ec0ed0,
    completionTag=0x7ffde1121260 "") at pquery.c:796
#7  0x00000000008de0d5 in exec_simple_query (query_string=0x2ec0068
"load 'user_acl';") at postgres.c:1215
#8  0x00000000008e2256 in PostgresMain (argc=1, argv=0x2ef4710,
dbname=0x2ebcd08 "postgres", username=0x2ef4558 "admin") at
postgres.c:4247
#9  0x0000000000839272 in BackendRun (port=0x2ee7430) at
postmaster.c:4448
#10 0x0000000000838a38 in BackendStartup (port=0x2ee7430) at
postmaster.c:4139
#11 0x0000000000834df1 in ServerLoop () at postmaster.c:1704
#12 0x00000000008346c0 in PostmasterMain (argc=3, argv=0x2ebac80) at
postmaster.c:1377
#13 0x0000000000755bd8 in main (argc=3, argv=0x2ebac80) at main.c:228
(gdb)
```

功能描述

- 卸载插件时，负责插件的清理工作。
- 源文件 `src/backend/utils/fmgr/dfmgr.c` 中的函数 `internal_unload_library` 通过以下代码调用插件的清理函数 `_PG_fini`。

```
/*
 * If the library has a _PG_fini() function, call it.
 */
PG_fini = (PG_fini_t) dlsym(file_scanner->handle, "_PG_fini");
if (PG_fini)
```

- 插件 `pg_stat_statements` 公共该函数恢复在初始化函数中修改过的钩子。

函数代码

```
/*
 * Module unload callback
 */
void
_PG_fini(void)
{
    /* Uninstall hooks. */
    shmem_startup_hook = prev_shmem_startup_hook;
    post_parse_analyze_hook = prev_post_parse_analyze_hook;
    ExecutorStart_hook = prev_ExecutorStart;
    ExecutorRun_hook = prev_ExecutorRun;
    ExecutorFinish_hook = prev_ExecutorFinish;
    ExecutorEnd_hook = prev_ExecutorEnd;
    ProcessUtility_hook = prev_ProcessUtility;
}
```

函数 `pgss_shmem_startup`-创建共享内存

函数代码

```

/*
 * shmem_startup hook: allocate or attach to shared memory,
 * then load any pre-existing statistics from file.
 * Also create and load the query-texts file, which is expected to exist
 * (even if empty) while the module is enabled.
 */
static void
pgss_shmem_startup(void)
{
    bool    found;
    HASHCTL  info;
    FILE     *file = NULL;
    FILE     *qfile = NULL;
    uint32   header;
    int32    num;
    int32    pgver;
    int32    i;
    int      buffer_size;
    char     *buffer = NULL;

    if (prev_shmem_startup_hook)
        prev_shmem_startup_hook();

    /* reset in case this is a restart within the postmaster */
    pgss = NULL;
    pgss_hash = NULL;

    /*
     * Create or attach to the shared memory state, including hash table
     */
    LWLockAcquire(AddinShmemInitLock, LW_EXCLUSIVE);

    pgss = ShmemInitStruct("pg_stat_statements",
                          sizeof(pgssSharedState),
                          &found);

    if (!found)
    {
        /* First time through ... */
        pgss->lock = &(GetNamedLWLockTranche("pg_stat_statements"))->lock;
        pgss->cur_median_usage = ASSUMED_MEDIAN_INIT;
        pgss->mean_query_len = ASSUMED_LENGTH_INIT;
        SpinLockInit(&pgss->mutex);
        pgss->extent = 0;
        pgss->n_writers = 0;
        pgss->gc_count = 0;
    }
}

```



```

memset(&info, 0, sizeof(info));
info.keysize = sizeof(pgssHashKey);
info.entrysize = sizeof(pgssEntry);
pgss_hash = ShmemInitHash("pg_stat_statements hash",
                          pgss_max, pgss_max,
                          &info,
                          HASH_ELEM | HASH_BLOBS);

LWLockRelease(AddinShmemInitLock);

/*
 * If we're in the postmaster (or a standalone backend...), set up a shmem
 * exit hook to dump the statistics to disk.
 */
if (!IsUnderPostmaster)
    on_shmem_exit(pgss_shmem_shutdown, (Datum) 0);

/*
 * Done if some other process already completed our initialization.
 */
if (found)
    return;

/*
 * Note: we don't bother with locks here, because there should be no other
 * processes running when this code is reached.
 */

/* Unlink query text file possibly left over from crash */
unlink(PGSS_TEXT_FILE);

/* Allocate new query text temp file */
qfile = AllocateFile(PGSS_TEXT_FILE, PG_BINARY_W);
if (qfile == NULL)
    goto write_error;

/*
 * If we were told not to load old statistics, we're done. (Note we do
 * not try to unlink any old dump file in this case. This seems a bit
 * questionable but it's the historical behavior.)
 */
if (!pgss_save)
{
    FreeFile(qfile);
    return;
}

/*
 * Attempt to load old statistics from the dump file.

```

```

*/
file = AllocateFile(PGSS_DUMP_FILE, PG_BINARY_R);
if (file == NULL)
{
    if (errno != ENOENT)
        goto read_error;
    /* No existing persisted stats file, so we're done */
    FreeFile(qfile);
    return;
}

buffer_size = 2048;
buffer = (char *) palloc(buffer_size);

if (fread(&header, sizeof(uint32), 1, file) != 1 ||
    fread(&pgver, sizeof(uint32), 1, file) != 1 ||
    fread(&num, sizeof(int32), 1, file) != 1)
    goto read_error;

if (header != PGSS_FILE_HEADER ||
    pgver != PGSS_PG_MAJOR_VERSION)
    goto data_error;

for (i = 0; i < num; i++)
{
    pgssEntry temp;
    pgssEntry *entry;
    Size    query_offset;

    if (fread(&temp, sizeof(pgssEntry), 1, file) != 1)
        goto read_error;

    /* Encoding is the only field we can easily sanity-check */
    if (!PG_VALID_BE_ENCODING(temp.encoding))
        goto data_error;

    /* Resize buffer as needed */
    if (temp.query_len >= buffer_size)
    {
        buffer_size = Max(buffer_size * 2, temp.query_len + 1);
        buffer = repalloc(buffer, buffer_size);
    }

    if (fread(buffer, 1, temp.query_len + 1, file) != temp.query_len + 1)
        goto read_error;

    /* Should have a trailing null, but let's make sure */
    buffer[temp.query_len] = '\0';
}

```

```

/* Skip loading "sticky" entries */
if (temp.counters.calls == 0)
    continue;

/* Store the query text */
query_offset = pgss->extent;
if (fwrite(buffer, 1, temp.query_len + 1, qfile) != temp.query_len + 1)
    goto write_error;
pgss->extent += temp.query_len + 1;

/* make the hashtable entry (discards old entries if too many) */
entry = entry_alloc(&temp.key, query_offset, temp.query_len,
                    temp.encoding,
                    false);

/* copy in the actual stats */
entry->counters = temp.counters;
}

pfree(buffer);
FreeFile(file);
FreeFile(qfile);

/*
 * Remove the persisted stats file so it's not included in
 * backups/replication standbys, etc. A new file will be written on next
 * shutdown.
 *
 * Note: it's okay if the PGSS_TEXT_FILE is included in a basebackup,
 * because we remove that file on startup; it acts inversely to
 * PGSS_DUMP_FILE, in that it is only supposed to be around when the
 * server is running, whereas PGSS_DUMP_FILE is only supposed to be around
 * when the server is not running. Leaving the file creates no danger of
 * a newly restored database having a spurious record of execution costs,
 * which is what we're really concerned about here.
 */
unlink(PGSS_DUMP_FILE);

return;

read_error:
ereport(LOG,
        (errcode_for_file_access(),
         errmsg("could not read file \"%s\": %m",
                PGSS_DUMP_FILE)));
goto fail;

data_error:
ereport(LOG,
        (errcode(ERRCODE_INVALID_PARAMETER_VALUE),

```

```

        errmsg("ignoring invalid data in file \"%s\"",
                PGSS_DUMP_FILE));
    goto fail;
write_error:
    ereport(LOG,
        (errcode_for_file_access(),
         errmsg("could not write file \"%s\": %m",
                PGSS_TEXT_FILE)));
fail:
    if (buffer)
        pfree(buffer);
    if (file)
        FreeFile(file);
    if (qfile)
        FreeFile(qfile);
    /* If possible, throw away the bogus file; ignore any error */
    unlink(PGSS_DUMP_FILE);

    /*
     * Don't unlink PGSS_TEXT_FILE here; it should always be around while the
     * server is running with pg_stat_statements enabled
     */
}

```

函数 `pgss_shmem_shutdown` - 将统计信息 dump 到文件

函数代码

```

/*
 * shmem_shutdown hook: Dump statistics into file.
 *
 * Note: we don't bother with acquiring lock, because there should be no
 * other processes running when this is called.
 */
static void
pgss_shmem_shutdown(int code, Datum arg)
{
    FILE      *file;
    char      *qbuffer = NULL;
    Size      qbuffer_size = 0;
    HASH_SEQ_STATUS hash_seq;

```

```

int32    num_entries;
pgssEntry *entry;

/* Don't try to dump during a crash. */
if (code)
    return;

/* Safety check ... shouldn't get here unless shmem is set up. */
if (!pgss || !pgss_hash)
    return;

/* Don't dump if told not to. */
if (!pgss_save)
    return;

file = AllocateFile(PGSS_DUMP_FILE ".tmp", PG_BINARY_W);
if (file == NULL)
    goto error;

if (fwrite(&PGSS_FILE_HEADER, sizeof(uint32), 1, file) != 1)
    goto error;
if (fwrite(&PGSS_PG_MAJOR_VERSION, sizeof(uint32), 1, file) != 1)
    goto error;
num_entries = hash_get_num_entries(pgss_hash);
if (fwrite(&num_entries, sizeof(int32), 1, file) != 1)
    goto error;

qbuffer = qtext_load_file(&qbuffer_size);
if (qbuffer == NULL)
    goto error;

/*
 * When serializing to disk, we store query texts immediately after their
 * entry data. Any orphaned query texts are thereby excluded.
 */
hash_seq_init(&hash_seq, pgss_hash);
while ((entry = hash_seq_search(&hash_seq)) != NULL)
{
    int    len = entry->query_len;
    char    *qstr = qtext_fetch(entry->query_offset, len,
                                qbuffer, qbuffer_size);

    if (qstr == NULL)
        continue;    /* Ignore any entries with bogus texts */

    if (fwrite(entry, sizeof(pgssEntry), 1, file) != 1 ||
        fwrite(qstr, 1, len + 1, file) != len + 1)
    {
        /* note: we assume hash_seq_term won't change errno */

```

```

        hash_seq_term(&hash_seq);
        goto error;
    }
}

free(qbuffer);
qbuffer = NULL;

if (FreeFile(file))
{
    file = NULL;
    goto error;
}

/*
 * Rename file into place, so we atomically replace any old one.
 */
(void) durable_rename(PGSS_DUMP_FILE ".tmp", PGSS_DUMP_FILE, LOG);

/* Unlink query-texts file; it's not needed while shutdown */
unlink(PGSS_TEXT_FILE);

return;

error:
    ereport(LOG,
        (errcode_for_file_access(),
         errmsg("could not write file \"%s\": %m",
                PGSS_DUMP_FILE ".tmp")));
    if (qbuffer)
        free(qbuffer);
    if (file)
        FreeFile(file);
    unlink(PGSS_DUMP_FILE ".tmp");
    unlink(PGSS_TEXT_FILE);
}

```

函数 pgss_post_parse_analyze

函数代码

```

/*
 * Post-parse-analysis hook: mark query with a queryId

```

```

*/
static void
pgss_post_parse_analyze(ParseState *pstate, Query *query)
{
    pgssJumbleState jstate;

    if (prev_post_parse_analyze_hook)
        prev_post_parse_analyze_hook(pstate, query);

    /* Assert we didn't do this already */
    Assert(query->queryId == UINT64CONST(0));

    /* Safety check... */
    if (!pgss || !pgss_hash)
        return;

    /*
     * Utility statements get queryId zero. We do this even in cases where
     * the statement contains an optimizable statement for which a queryId
     * could be derived (such as EXPLAIN or DECLARE CURSOR). For such cases,
     * runtime control will first go through ProcessUtility and then the
     * executor, and we don't want the executor hooks to do anything, since we
     * are already measuring the statement's costs at the utility level.
     */
    if (query->utilityStmt)
    {
        query->queryId = UINT64CONST(0);
        return;
    }

    /* Set up workspace for query jumbling */
    jstate.jumble = (unsigned char *) palloc(JUMBLE_SIZE);
    jstate.jumble_len = 0;
    jstate.clocations_buf_size = 32;
    jstate.clocations = (pgssLocationLen *)
        palloc(jstate.clocations_buf_size * sizeof(pgssLocationLen));
    jstate.clocations_count = 0;
    jstate.highest_extern_param_id = 0;

    /* Compute query ID and mark the Query node with it */
    JumbleQuery(&jstate, query);
    query->queryId =
        DatumGetUInt64(hash_any_extended(jstate.jumble, jstate.jumble_len, 0));

    /*
     * If we are unlucky enough to get a hash of zero, use 1 instead, to
     * prevent confusion with the utility-statement case.
     */
    if (query->queryId == UINT64CONST(0))

```

```

    query->queryId = UINT64CONST(1);

/*
 * If we were able to identify any ignorable constants, we immediately
 * create a hash table entry for the query, so that we can record the
 * normalized form of the query string.  If there were no such constants,
 * the normalized string would be the same as the query text anyway, so
 * there's no need for an early entry.
 */
if (jstate.clocations_count > 0)
    pgss_store(pstate->p_sourcetext,
               query->queryId,
               query->stmt_location,
               query->stmt_len,
               0,
               0,
               NULL,
               &jstate);
}

```

函数 `pgss_ExecutorStart` - 开启sql语句统计跟踪

功能描述

- 在执行完函数 `standard_ExecutorStart` 后，初始化与当前sql语句相关的统计信息，如下：

```

/*
 * Set up to track total elapsed time in ExecutorRun.  Make sure the
 * space is allocated in the per-query context so it will go away at
 * ExecutorEnd.
 */
if (queryDesc->totaltime == NULL)
{
    MemoryContext oldcxt;

    oldcxt = MemoryContextSwitchTo(queryDesc->estate->es_query_cxt);
    queryDesc->totaltime = InstrAlloc(1, INSTRUMENT_ALL);
    MemoryContextSwitchTo(oldcxt);
}

```

函数代码


```

/*
 * ExecutorStart hook: start up tracking if needed
 */
static void
pgss_ExecutorStart(QueryDesc *queryDesc, int eflags)
{
    if (prev_ExecutorStart)
        prev_ExecutorStart(queryDesc, eflags);
    else
        standard_ExecutorStart(queryDesc, eflags);

    /*
     * If query has queryId zero, don't track it. This prevents double
     * counting of optimizable statements that are directly contained in
     * utility statements.
     */
    if (pgss_enabled() && queryDesc->plannedstmt->queryId != UINT64CONST(0))
    {
        /*
         * Set up to track total elapsed time in ExecutorRun. Make sure the
         * space is allocated in the per-query context so it will go away at
         * ExecutorEnd.
         */
        if (queryDesc->totaltime == NULL)
        {
            MemoryContext oldcxt;

            oldcxt = MemoryContextSwitchTo(queryDesc->estate->es_query_cxt);
            queryDesc->totaltime = InstrAlloc(1, INSTRUMENT_ALL);
            MemoryContextSwitchTo(oldcxt);
        }
    }
}

```

函数 pgss_ProcessUtility

功能描述

- 工具命令是除了 `SELECT`、`INSERT`、`UPDATE` 和 `DELETE` 之外所有的其他命令。
- 该函数负责处理工具命令的统计信息。

钩子的定义

- 源文件

```
src/backend/tcop/utility.c
```

- 钩子的定义

```
/* Hook for plugins to get control in ProcessUtility() */  
ProcessUtility_hook_type ProcessUtility_hook = NULL;
```

- 钩子的函数原型

```
src/include/tcop/utility.h
```

```
/* Hook for plugins to get control in ProcessUtility() */  
typedef void (*ProcessUtility_hook_type) (PlannedStmt *pstmt,  
                                           const char *queryString, ProcessUtilityContext  
context,  
                                           ParamListInfo params,  
                                           QueryEnvironment *queryEnv,  
                                           DestReceiver *dest, char *completionTag);
```

函数代码

```
/*  
 * ProcessUtility hook  
 */  
static void  
pgss_ProcessUtility(PlannedStmt *pstmt, const char *queryString,  
                    ProcessUtilityContext context,  
                    ParamListInfo params, QueryEnvironment *queryEnv,  
                    DestReceiver *dest, char *completionTag)  
{  
    Node *parsetree = pstmt->utilityStmt;  
  
    /*  
     * If it's an EXECUTE statement, we don't track it and don't increment the  
     * nesting level. This allows the cycles to be charged to the underlying  
     * PREPARE instead (by the Executor hooks), which is much more useful.  
     */  
  
    /* We also don't track execution of PREPARE. If we did, we would get one  
     * hash table entry for the PREPARE (with hash calculated from the query
```

```

* string), and then a different one with the same query string (but hash
* calculated from the query tree) would be used to accumulate costs of
* ensuing EXECUTES. This would be confusing, and inconsistent with other
* cases where planning time is not included at all.
*
* Likewise, we don't track execution of DEALLOCATE.
*/
if (pgss_track_utility && pgss_enabled() &&
    !IsA(parsetree, ExecuteStmt) &&
    !IsA(parsetree, PrepareStmt) &&
    !IsA(parsetree, DeallocateStmt))
{
    instr_time  start;
    instr_time  duration;
    uint64      rows;
    BufferUsage  bufusage_start,
                bufusage;

    bufusage_start = pgBufferUsage;
    INSTR_TIME_SET_CURRENT(start);

    nested_level++;
    PG_TRY();
    {
        if (prev_ProcessUtility)
            prev_ProcessUtility(pstmt, queryString,
                                context, params, queryEnv,
                                dest, completionTag);
        else
            standard_ProcessUtility(pstmt, queryString,
                                    context, params, queryEnv,
                                    dest, completionTag);
        nested_level--;
    }
    PG_CATCH();
    {
        nested_level--;
        PG_RE_THROW();
    }
    PG_END_TRY();

    INSTR_TIME_SET_CURRENT(duration);
    INSTR_TIME_SUBTRACT(duration, start);

    /* parse command tag to retrieve the number of affected rows. */
    if (completionTag &&
        strncmp(completionTag, "COPY ", 5) == 0)
        rows = pg_strtoint64(completionTag + 5, NULL, 10);
    else

```

```

    rows = 0;

    /* calc differences of buffer counters. */
    bufusage.shared_blks_hit =
        pgBufferUsage.shared_blks_hit - bufusage_start.shared_blks_hit;
    bufusage.shared_blks_read =
        pgBufferUsage.shared_blks_read - bufusage_start.shared_blks_read;
    bufusage.shared_blks_dirtied =
        pgBufferUsage.shared_blks_dirtied - bufusage_start.shared_blks_dirtied;
    bufusage.shared_blks_written =
        pgBufferUsage.shared_blks_written - bufusage_start.shared_blks_written;
    bufusage.local_blks_hit =
        pgBufferUsage.local_blks_hit - bufusage_start.local_blks_hit;
    bufusage.local_blks_read =
        pgBufferUsage.local_blks_read - bufusage_start.local_blks_read;
    bufusage.local_blks_dirtied =
        pgBufferUsage.local_blks_dirtied - bufusage_start.local_blks_dirtied;
    bufusage.local_blks_written =
        pgBufferUsage.local_blks_written - bufusage_start.local_blks_written;
    bufusage.temp_blks_read =
        pgBufferUsage.temp_blks_read - bufusage_start.temp_blks_read;
    bufusage.temp_blks_written =
        pgBufferUsage.temp_blks_written - bufusage_start.temp_blks_written;
    bufusage.blk_read_time = pgBufferUsage.blk_read_time;
    INSTR_TIME_SUBTRACT(bufusage.blk_read_time, bufusage_start.blk_read_time);
    bufusage.blk_write_time = pgBufferUsage.blk_write_time;
    INSTR_TIME_SUBTRACT(bufusage.blk_write_time,
bufusage_start.blk_write_time);

    pgss_store(queryString,
        0,      /* signal that it's a utility stmt */
        pstmt->stmt_location,
        pstmt->stmt_len,
        INSTR_TIME_GET_MILLISEC(duration),
        rows,
        &bufusage,
        NULL);
}
else
{
    if (prev_ProcessUtility)
        prev_ProcessUtility(pstmt, queryString,
            context, params, queryEnv,
            dest, completionTag);
    else
        standard_ProcessUtility(pstmt, queryString,
            context, params, queryEnv,
            dest, completionTag);
}

```

```
}
```

与工具命令相关的其他函数

函数 `ProcessUtility`-工具命令处理函数

- 源文件

```
src/backend/tcop/utility.c
```

- 函数工具

- 工具命令的总入口，会调用函数 `standard_ProcessUtility` 执行具体的操作
- 调用钩子函数 `ProcessUtility_hook`

- 函数代码

```
/*
 * ProcessUtility
 *     general utility function invoker
 *
 * pstmt: PlannedStmt wrapper for the utility statement
 * queryString: original source text of command
 * context: identifies source of statement (toplevel client command,
 *         non-toplevel client command, subcommand of a larger utility command)
 * params: parameters to use during execution
 * queryEnv: environment for parse through execution (e.g., ephemeral named
 *         tables like trigger transition tables). May be NULL.
 * dest: where to send results
 * completionTag: points to a buffer of size COMPLETION_TAG_BUFSIZE
 *         in which to store a command completion status string.
 *
 * Caller MUST supply a queryString; it is not allowed (anymore) to pass NULL.
 * If you really don't have source text, you can pass a constant string,
 * perhaps "(query not available)".
 *
 * completionTag is only set nonempty if we want to return a nondefault status.
 *
 * completionTag may be NULL if caller doesn't want a status string.
 *
 * Note for users of ProcessUtility_hook: the same queryString may be passed
 * to multiple invocations of ProcessUtility when processing a query string
 * containing multiple semicolon-separated statements. One should use
 * pstmt->stmt_location and pstmt->stmt_len to identify the substring
 * containing the current statement. Keep in mind also that some utility
```

```

* statements (e.g., CREATE SCHEMA) will recurse to ProcessUtility to process
* sub-statements, often passing down the same queryString, stmt_location,
* and stmt_len that were given for the whole statement.
*/
void
ProcessUtility(PlannedStmt *pstmt,
               const char *queryString,
               ProcessUtilityContext context,
               ParamListInfo params,
               QueryEnvironment *queryEnv,
               DestReceiver *dest,
               char *completionTag)
{
    Assert(IsA(pstmt, PlannedStmt));
    Assert(pstmt->commandType == CMD_UTILITY);
    Assert(queryString != NULL); /* required as of 8.4 */

    /*
     * We provide a function hook variable that lets loadable plugins get
     * control when ProcessUtility is called. Such a plugin would normally
     * call standard_ProcessUtility().
     */
    if (ProcessUtility_hook)
        (*ProcessUtility_hook) (pstmt, queryString,
                                context, params, queryEnv,
                                dest, completionTag);
    else
        standard_ProcessUtility(pstmt, queryString,
                                context, params, queryEnv,
                                dest, completionTag);
}

```

函数 `standard_ProcessUtility`-工具命令处理函数

源文件

```
src/backend/tcop/utility.c
```

函数代码

```

/*
 * standard_ProcessUtility itself deals only with utility commands for
 * which we do not provide event trigger support. Commands that do have
 * such support are passed down to ProcessUtilitySlow, which contains the
 * necessary infrastructure for such triggers.
 *
 * This division is not just for performance: it's critical that the
 * event trigger code not be invoked when doing START TRANSACTION for
 * example, because we might need to refresh the event trigger cache,
 * which requires being in a valid transaction.
 */
void
standard_ProcessUtility(PlannedStmt *pstmt,
                        const char *queryString,
                        ProcessUtilityContext context,
                        ParamListInfo params,
                        QueryEnvironment *queryEnv,
                        DestReceiver *dest,
                        char *completionTag)
{
    Node      *parsetree = pstmt->utilityStmt;
    bool      isTopLevel = (context == PROCESS_UTILITY_TOPLEVEL);
    bool      isAtomicContext = (!(context == PROCESS_UTILITY_TOPLEVEL || context
== PROCESS_UTILITY_QUERY_NONATOMIC) || IsTransactionBlock());
    ParseState *pstate;

    /* This can recurse, so check for excessive recursion */
    check_stack_depth();

    check_xact_readonly(parsetree);

    if (completionTag)
        completionTag[0] = '\0';

    pstate = make_parsestate(NULL);
    pstate->p_sourcetext = queryString;

    switch (nodeTag(parsetree))
    {
        /*
         * ***** transactions *****
         */
        case T_TransactionStmt:
        {
            TransactionStmt *stmt = (TransactionStmt *) parsetree;

            switch (stmt->kind)
            {

```

```

/*
 * START TRANSACTION, as defined by SQL99: Identical
 * to BEGIN. Same code for both.
 */
case TRANS_STMT_BEGIN:
case TRANS_STMT_START:
{
    ListCell    *lc;

    BeginTransactionBlock();
    foreach(lc, stmt->options)
    {
        DefElem    *item = (DefElem *) lfirst(lc);

        if (strcmp(item->defname, "transaction_isolation") == 0)
            SetPGVariable("transaction_isolation",
                          list_makel(item->arg),
                          true);
        else if (strcmp(item->defname, "transaction_read_only") == 0)
            SetPGVariable("transaction_read_only",
                          list_makel(item->arg),
                          true);
        else if (strcmp(item->defname, "transaction_deferrable") == 0)
            SetPGVariable("transaction_deferrable",
                          list_makel(item->arg),
                          true);
    }
}
break;

case TRANS_STMT_COMMIT:
    if (!EndTransactionBlock(stmt->chain))
    {
        /* report unsuccessful commit in completionTag */
        if (completionTag)
            strcpy(completionTag, "ROLLBACK");
    }
    break;

case TRANS_STMT_PREPARE:
    PreventCommandDuringRecovery("PREPARE TRANSACTION");
    if (!PrepareTransactionBlock(stmt->gid))
    {
        /* report unsuccessful commit in completionTag */
        if (completionTag)
            strcpy(completionTag, "ROLLBACK");
    }
    break;

```



```

    case TRANS_STMT_COMMIT_PREPARED:
        PreventInTransactionBlock(isTopLevel, "COMMIT PREPARED");
        PreventCommandDuringRecovery("COMMIT PREPARED");
        FinishPreparedTransaction(stmt->gid, true);
        break;

    case TRANS_STMT_ROLLBACK_PREPARED:
        PreventInTransactionBlock(isTopLevel, "ROLLBACK PREPARED");
        PreventCommandDuringRecovery("ROLLBACK PREPARED");
        FinishPreparedTransaction(stmt->gid, false);
        break;

    case TRANS_STMT_ROLLBACK:
        UserAbortTransactionBlock(stmt->chain);
        break;

    case TRANS_STMT_SAVEPOINT:
        RequireTransactionBlock(isTopLevel, "SAVEPOINT");
        DefineSavepoint(stmt->savepoint_name);
        break;

    case TRANS_STMT_RELEASE:
        RequireTransactionBlock(isTopLevel, "RELEASE SAVEPOINT");
        ReleaseSavepoint(stmt->savepoint_name);
        break;

    case TRANS_STMT_ROLLBACK_TO:
        RequireTransactionBlock(isTopLevel, "ROLLBACK TO SAVEPOINT");
        RollbackToSavepoint(stmt->savepoint_name);

        /*
         * CommitTransactionCommand is in charge of
         * re-defining the savepoint again
         */
        break;
}
}
break;

/*
 * Portal (cursor) manipulation
 */
case T_DeclareCursorStmt:
    PerformCursorOpen((DeclareCursorStmt *) parsetree, params,
                      queryString, isTopLevel);
    break;

case T_ClosePortalStmt:
{

```

```

        ClosePortalStmt *stmt = (ClosePortalStmt *) parsetree;

        CheckRestrictedOperation("CLOSE");
        PerformPortalClose(stmt->portalname);
    }
    break;

case T_FetchStmt:
    PerformPortalFetch((FetchStmt *) parsetree, dest,
                      completionTag);
    break;

case T_DoStmt:
    ExecuteDoStmt((DoStmt *) parsetree, isAtomicContext);
    break;

case T_CreateTableSpaceStmt:
    /* no event triggers for global objects */
    PreventInTransactionBlock(isTopLevel, "CREATE TABLESPACE");
    CreateTableSpace((CreateTableSpaceStmt *) parsetree);
    break;

case T_DropTableSpaceStmt:
    /* no event triggers for global objects */
    PreventInTransactionBlock(isTopLevel, "DROP TABLESPACE");
    DropTableSpace((DropTableSpaceStmt *) parsetree);
    break;

case T_AlterTableSpaceOptionsStmt:
    /* no event triggers for global objects */
    AlterTableSpaceOptions((AlterTableSpaceOptionsStmt *) parsetree);
    break;

case T_TruncateStmt:
    ExecuteTruncate((TruncateStmt *) parsetree);
    break;

case T_CopyStmt:
    {
        uint64    processed;

        DoCopy(pstate, (CopyStmt *) parsetree,
              pstmt->stmt_location, pstmt->stmt_len,
              &processed);
        if (completionTag)
            snprintf(completionTag, COMPLETION_TAG_BUFSIZE,
                     "COPY " UINT64_FORMAT, processed);
    }
    break;

```

```

case T_PrepareStmt:
    CheckRestrictedOperation("PREPARE");
    PrepareQuery((PrepareStmt *) parsetree, queryString,
        pstmt->stmt_location, pstmt->stmt_len);
    break;

case T_ExecuteStmt:
    ExecuteQuery((ExecuteStmt *) parsetree, NULL,
        queryString, params,
        dest, completionTag);
    break;

case T_DeallocateStmt:
    CheckRestrictedOperation("DEALLOCATE");
    DeallocateQuery((DeallocateStmt *) parsetree);
    break;

case T_GrantRoleStmt:
    /* no event triggers for global objects */
    GrantRole((GrantRoleStmt *) parsetree);
    break;

case T_CreatedbStmt:
    /* no event triggers for global objects */
    PreventInTransactionBlock(isTopLevel, "CREATE DATABASE");
    createdb(pstate, (CreatedbStmt *) parsetree);
    break;

case T_AlterDatabaseStmt:
    /* no event triggers for global objects */
    AlterDatabase(pstate, (AlterDatabaseStmt *) parsetree, isTopLevel);
    break;

case T_AlterDatabaseSetStmt:
    /* no event triggers for global objects */
    AlterDatabaseSet((AlterDatabaseSetStmt *) parsetree);
    break;

case T_DropdbStmt:
    {
        DropdbStmt *stmt = (DropdbStmt *) parsetree;

        /* no event triggers for global objects */
        PreventInTransactionBlock(isTopLevel, "DROP DATABASE");
        dropdb(stmt->dbname, stmt->missing_ok);
    }
    break;

```

```

/* Query-level asynchronous notification */
case T_NotifyStmt:
{
    NotifyStmt *stmt = (NotifyStmt *) parsetree;

    PreventCommandDuringRecovery("NOTIFY");
    Async_Notify(stmt->conditionname, stmt->payload);
}
break;

case T_ListenStmt:
{
    ListenStmt *stmt = (ListenStmt *) parsetree;

    PreventCommandDuringRecovery("LISTEN");
    CheckRestrictedOperation("LISTEN");
    Async_Listen(stmt->conditionname);
}
break;

case T_UnlistenStmt:
{
    UnlistenStmt *stmt = (UnlistenStmt *) parsetree;

    /* we allow UNLISTEN during recovery, as it's a noop */
    CheckRestrictedOperation("UNLISTEN");
    if (stmt->conditionname)
        Async_Unlisten(stmt->conditionname);
    else
        Async_UnlistenAll();
}
break;

case T_LoadStmt:
{
    LoadStmt *stmt = (LoadStmt *) parsetree;

    closeAllVfds(); /* probably not necessary... */
    /* Allowed names are restricted if you're not superuser */
    load_file(stmt->filename, !superuser());
}
break;

case T_CallStmt:
    ExecuteCallStmt(castNode(CallStmt, parsetree), params, isAtomicContext,
dest);
    break;

case T_ClusterStmt:

```

```

/* we choose to allow this during "read only" transactions */
PreventCommandDuringRecovery("CLUSTER");
/* forbidden in parallel mode due to CommandIsReadOnly */
cluster((ClusterStmt *) parsetree, isTopLevel);
break;

case T_VacuumStmt:
{
    VacuumStmt *stmt = (VacuumStmt *) parsetree;

    /* we choose to allow this during "read only" transactions */
    PreventCommandDuringRecovery(stmt->is_vacuumcmd ?
        "VACUUM" : "ANALYZE");
    /* forbidden in parallel mode due to CommandIsReadOnly */
    ExecVacuum(pstate, stmt, isTopLevel);
}
break;

case T_ExplainStmt:
    ExplainQuery(pstate, (ExplainStmt *) parsetree, queryString, params,
        queryEnv, dest);
    break;

case T_AlterSystemStmt:
    PreventInTransactionBlock(isTopLevel, "ALTER SYSTEM");
    AlterSystemSetConfigFile((AlterSystemStmt *) parsetree);
    break;

case T_VariableSetStmt:
    ExecSetVariableStmt((VariableSetStmt *) parsetree, isTopLevel);
    break;

case T_VariableShowStmt:
{
    VariableShowStmt *n = (VariableShowStmt *) parsetree;

    GetPGVariable(n->name, dest);
}
break;

case T_DiscardStmt:
    /* should we allow DISCARD PLANS? */
    CheckRestrictedOperation("DISCARD");
    DiscardCommand((DiscardStmt *) parsetree, isTopLevel);
    break;

case T_CreateEventTrigStmt:
    /* no event triggers on event triggers */
    CreateEventTrigger((CreateEventTrigStmt *) parsetree);

```

```

        break;

case T_AlterEventTrigStmt:
    /* no event triggers on event triggers */
    AlterEventTrigger((AlterEventTrigStmt *) parsetree);
    break;

    /*
     * ***** ROLE statements *****
     */

case T_CreateRoleStmt:
    /* no event triggers for global objects */
    CreateRole(pstate, (CreateRoleStmt *) parsetree);
    break;

case T_AlterRoleStmt:
    /* no event triggers for global objects */
    AlterRole((AlterRoleStmt *) parsetree);
    break;

case T_AlterRoleSetStmt:
    /* no event triggers for global objects */
    AlterRoleSet((AlterRoleSetStmt *) parsetree);
    break;

case T_DropRoleStmt:
    /* no event triggers for global objects */
    DropRole((DropRoleStmt *) parsetree);
    break;

case T_ReassignOwnedStmt:
    /* no event triggers for global objects */
    ReassignOwnedObjects((ReassignOwnedStmt *) parsetree);
    break;

case T_LockStmt:

    /*
     * Since the lock would just get dropped immediately, LOCK TABLE
     * outside a transaction block is presumed to be user error.
     */
    RequireTransactionBlock(isTopLevel, "LOCK TABLE");
    /* forbidden in parallel mode due to CommandIsReadOnly */
    LockTableCommand((LockStmt *) parsetree);
    break;

case T_ConstraintsSetStmt:
    WarnNoTransactionBlock(isTopLevel, "SET CONSTRAINTS");
    AfterTriggerSetState((ConstraintsSetStmt *) parsetree);

```

```

break;

case T_CheckPointStmt:
    if (!superuser())
        ereport(ERROR,
            (errcode(ERRCODE_INSUFFICIENT_PRIVILEGE),
             errmsg("must be superuser to do CHECKPOINT")));

    /*
     * You might think we should have a PreventCommandDuringRecovery()
     * here, but we interpret a CHECKPOINT command during recovery as
     * a request for a restartpoint instead. We allow this since it
     * can be a useful way of reducing switchover time when using
     * various forms of replication.
     */
    RequestCheckpoint(CHECKPOINT_IMMEDIATE | CHECKPOINT_WAIT |
        (RecoveryInProgress() ? 0 : CHECKPOINT_FORCE));
    break;

case T_ReindexStmt:
{
    ReindexStmt *stmt = (ReindexStmt *) parsetree;

    if (stmt->concurrent)
        PreventInTransactionBlock(isTopLevel,
            "REINDEX CONCURRENTLY");

    /* we choose to allow this during "read only" transactions */
    PreventCommandDuringRecovery("REINDEX");
    /* forbidden in parallel mode due to CommandIsReadOnly */
    switch (stmt->kind)
    {
        case REINDEX_OBJECT_INDEX:
            ReindexIndex(stmt->relation, stmt->options, stmt->concurrent);
            break;
        case REINDEX_OBJECT_TABLE:
            ReindexTable(stmt->relation, stmt->options, stmt->concurrent);
            break;
        case REINDEX_OBJECT_SCHEMA:
        case REINDEX_OBJECT_SYSTEM:
        case REINDEX_OBJECT_DATABASE:

            /*
             * This cannot run inside a user transaction block; if
             * we were inside a transaction, then its commit- and
             * start-transaction-command calls would not have the
             * intended effect!
             */
            PreventInTransactionBlock(isTopLevel,

```

```

                                (stmt->kind == REINDEX_OBJECT_SCHEMA) ? "REINDEX
SCHEMA" :

                                (stmt->kind == REINDEX_OBJECT_SYSTEM) ? "REINDEX
SYSTEM" :

                                "REINDEX DATABASE");
    ReindexMultipleTables(stmt->name, stmt->kind, stmt->options, stmt-
>concurrent);
    break;
default:
    elog(ERROR, "unrecognized object type: %d",
         (int) stmt->kind);
    break;
}
}
break;

/*
 * The following statements are supported by Event Triggers only
 * in some cases, so we "fast path" them in the other cases.
 */

case T_GrantStmt:
{
    GrantStmt  *stmt = (GrantStmt *) parsetree;

    if (EventTriggerSupportsObjectType(stmt->objtype))
        ProcessUtilitySlow(pstate, pstmt, queryString,
                           context, params, queryEnv,
                           dest, completionTag);
    else
        ExecuteGrantStmt(stmt);
}
break;

case T_DropStmt:
{
    DropStmt   *stmt = (DropStmt *) parsetree;

    if (EventTriggerSupportsObjectType(stmt->removeType))
        ProcessUtilitySlow(pstate, pstmt, queryString,
                           context, params, queryEnv,
                           dest, completionTag);
    else
        ExecDropStmt(stmt, isTopLevel);
}
break;

case T_RenameStmt:
{

```



```

RenameStmt *stmt = (RenameStmt *) parsetree;

if (EventTriggerSupportsObjectType(stmt->renameType))
    ProcessUtilitySlow(pstate, pstmt, queryString,
                        context, params, queryEnv,
                        dest, completionTag);
else
    ExecRenameStmt(stmt);
}
break;

case T_AlterObjectDependsStmt:
{
    AlterObjectDependsStmt *stmt = (AlterObjectDependsStmt *) parsetree;

    if (EventTriggerSupportsObjectType(stmt->objectType))
        ProcessUtilitySlow(pstate, pstmt, queryString,
                            context, params, queryEnv,
                            dest, completionTag);
    else
        ExecAlterObjectDependsStmt(stmt, NULL);
}
break;

case T_AlterObjectSchemaStmt:
{
    AlterObjectSchemaStmt *stmt = (AlterObjectSchemaStmt *) parsetree;

    if (EventTriggerSupportsObjectType(stmt->objectType))
        ProcessUtilitySlow(pstate, pstmt, queryString,
                            context, params, queryEnv,
                            dest, completionTag);
    else
        ExecAlterObjectSchemaStmt(stmt, NULL);
}
break;

case T_AlterOwnerStmt:
{
    AlterOwnerStmt *stmt = (AlterOwnerStmt *) parsetree;

    if (EventTriggerSupportsObjectType(stmt->objectType))
        ProcessUtilitySlow(pstate, pstmt, queryString,
                            context, params, queryEnv,
                            dest, completionTag);
    else
        ExecAlterOwnerStmt(stmt);
}
break;

```

```

case T_CommentStmt:
{
    CommentStmt *stmt = (CommentStmt *) parsetree;

    if (EventTriggerSupportsObjectType(stmt->objtype))
        ProcessUtilitySlow(pstate, pstmt, queryString,
                           context, params, queryEnv,
                           dest, completionTag);
    else
        CommentObject(stmt);
    break;
}

case T_SecLabelStmt:
{
    SecLabelStmt *stmt = (SecLabelStmt *) parsetree;

    if (EventTriggerSupportsObjectType(stmt->objtype))
        ProcessUtilitySlow(pstate, pstmt, queryString,
                           context, params, queryEnv,
                           dest, completionTag);
    else
        ExecSecLabelStmt(stmt);
    break;
}

default:
    /* All other statement types have event trigger support */
    ProcessUtilitySlow(pstate, pstmt, queryString,
                       context, params, queryEnv,
                       dest, completionTag);
    break;
}

free_parsestate(pstate);

/*
 * Make effects of commands visible, for instance so that
 * PreCommit_on_commit_actions() can see them (see for example bug
 * #15631).
 */
CommandCounterIncrement();
}

```

函数 ProcessUtilitySlow

源文件

```
src/backend/tcop/utility.c
```

函数代码

```
/*
 * The "Slow" variant of ProcessUtility should only receive statements
 * supported by the event triggers facility. Therefore, we always
 * perform the trigger support calls if the context allows it.
 */
static void
ProcessUtilitySlow(ParseState *pstate,
                   PlannedStmt *pstmt,
                   const char *queryString,
                   ProcessUtilityContext context,
                   ParamListInfo params,
                   QueryEnvironment *queryEnv,
                   DestReceiver *dest,
                   char *completionTag)
{
    Node *parsetree = pstmt->utilityStmt;
    bool isTopLevel = (context == PROCESS_UTILITY_TOPLEVEL);
    bool isCompleteQuery = (context != PROCESS_UTILITY_SUBCOMMAND);
    bool needCleanup;
    bool commandCollected = false;
    ObjectAddress address;
    ObjectAddress secondaryObject = InvalidObjectAddress;

    /* All event trigger calls are done only when isCompleteQuery is true */
    needCleanup = isCompleteQuery && EventTriggerBeginCompleteQuery();

    /* PG_TRY block is to ensure we call EventTriggerEndCompleteQuery */
    PG_TRY();
    {
        if (isCompleteQuery)
            EventTriggerDDLCommandStart(parsetree);

        switch (nodeTag(parsetree))
        {
            /*
             * relation and attribute manipulation
             */
            case T_CreateSchemaStmt:
                CreateSchemaCommand((CreateSchemaStmt *) parsetree,
```

```

        queryString,
        pstmt->stmt_location,
        pstmt->stmt_len);

/*
 * EventTriggerCollectSimpleCommand called by
 * CreateSchemaCommand
 */
commandCollected = true;
break;

case T_CreateStmt:
case T_CreateForeignTableStmt:
{
    List      *stmts;
    ListCell   *l;

    /* Run parse analysis ... */
    stmts = transformCreateStmt((CreateStmt *) parsetree,
                                queryString);

    /* ... and do it */
    foreach(l, stmts)
    {
        Node      *stmt = (Node *) lfirst(l);

        if (IsA(stmt, CreateStmt))
        {
            Datum   toast_options;
            static char *validnsp[] = HEAP_RELOPT_NAMESPACES;

            /* Create the table itself */
            address = DefineRelation((CreateStmt *) stmt,
                                    RELKIND_RELATION,
                                    InvalidOid, NULL,
                                    queryString);
            EventTriggerCollectSimpleCommand(address,
                                             secondaryObject,
                                             stmt);

            /*
             * Let NewRelationCreateToastTable decide if this
             * one needs a secondary relation too.
             */
            CommandCounterIncrement();

            /*
             * parse and validate reloptions for the toast
             * table

```

```

    */
    toast_options = transformRelOptions((Datum) 0,
                                        ((CreateStmt *) stmt)->options,
                                        "toast",
                                        validnsp,
                                        true,
                                        false);
    (void) heap_reloptions(RELKIND_TOASTVALUE,
                          toast_options,
                          true);

    NewRelationCreateToastTable(address.objectId,
                               toast_options);
}
else if (IsA(stmt, CreateForeignTableStmt))
{
    /* Create the table itself */
    address = DefineRelation((CreateStmt *) stmt,
                           RELKIND_FOREIGN_TABLE,
                           InvalidOid, NULL,
                           queryString);
    CreateForeignTable((CreateForeignTableStmt *) stmt,
                     address.objectId);
    EventTriggerCollectSimpleCommand(address,
                                     secondaryObject,
                                     stmt);
}
else
{
    /*
     * Recurse for anything else. Note the recursive
     * call will stash the objects so created into our
     * event trigger context.
     */
    PlannedStmt *wrapper;

    wrapper = makeNode(PlannedStmt);
    wrapper->commandType = CMD_UTILITY;
    wrapper->canSetTag = false;
    wrapper->utilityStmt = stmt;
    wrapper->stmt_location = pstmt->stmt_location;
    wrapper->stmt_len = pstmt->stmt_len;

    ProcessUtility(wrapper,
                  queryString,
                  PROCESS_UTILITY_SUBCOMMAND,
                  params,
                  NULL,
                  None_Receiver,

```

```

        NULL);
    }

    /* Need CCI between commands */
    if (lnext(l) != NULL)
        CommandCounterIncrement();
}

/*
 * The multiple commands generated here are stashed
 * individually, so disable collection below.
 */
commandCollected = true;
}
break;

case T_AlterTableStmt:
{
    AlterTableStmt *atstmt = (AlterTableStmt *) parsetree;
    Oid            relid;
    List           *stmts;
    ListCell       *l;
    LOCKMODE       lockmode;

    /*
     * Figure out lock mode, and acquire lock. This also does
     * basic permissions checks, so that we won't wait for a
     * lock on (for example) a relation on which we have no
     * permissions.
     */
    lockmode = AlterTableGetLockLevel(atstmt->cmds);
    relid = AlterTableLookupRelation(atstmt, lockmode);

    if (OidIsValid(relid))
    {
        /* Run parse analysis ... */
        stmts = transformAlterTableStmt(relid, atstmt,
                                         queryString);

        /* ... ensure we have an event trigger context ... */
        EventTriggerAlterTableStart(parsetree);
        EventTriggerAlterTableRelid(relid);

        /* ... and do it */
        foreach(l, stmts)
        {
            Node      *stmt = (Node *) lfirst(l);

            if (IsA(stmt, AlterTableStmt))

```

```

{
    /* Do the table alteration proper */
    AlterTable(relid, lockmode,
               (AlterTableStmt *) stmt);
}
else
{
    /*
     * Recurse for anything else.  If we need to
     * do so, "close" the current complex-command
     * set, and start a new one at the bottom;
     * this is needed to ensure the ordering of
     * queued commands is consistent with the way
     * they are executed here.
     */
    PlannedStmt *wrapper;

    EventTriggerAlterTableEnd();
    wrapper = makeNode(PlannedStmt);
    wrapper->commandType = CMD_UTILITY;
    wrapper->canSetTag = false;
    wrapper->utilityStmt = stmt;
    wrapper->stmt_location = pstmt->stmt_location;
    wrapper->stmt_len = pstmt->stmt_len;
    ProcessUtility(wrapper,
                   queryString,
                   PROCESS_UTILITY_SUBCOMMAND,
                   params,
                   NULL,
                   None_Receiver,
                   NULL);

    EventTriggerAlterTableStart(parsetree);
    EventTriggerAlterTableRelid(relid);
}

/* Need CCI between commands */
if (lnext(l) != NULL)
    CommandCounterIncrement();
}

/* done */
EventTriggerAlterTableEnd();
}
else
    ereport(NOTICE,
            (errmsg("relation \"%s\" does not exist, skipping",
                  atstmt->relation->relname)));
}

```

```

/* ALTER TABLE stashes commands internally */
commandCollected = true;
break;

case T_AlterDomainStmt:
{
    AlterDomainStmt *stmt = (AlterDomainStmt *) parsetree;

    /*
     * Some or all of these functions are recursive to cover
     * inherited things, so permission checks are done there.
     */
    switch (stmt->subtype)
    {
        case 'T': /* ALTER DOMAIN DEFAULT */

            /*
             * Recursively alter column default for table and,
             * if requested, for descendants
             */
            address =
                AlterDomainDefault(stmt->typeName,
                                   stmt->def);

            break;
        case 'N': /* ALTER DOMAIN DROP NOT NULL */
            address =
                AlterDomainNotNull(stmt->typeName,
                                   false);

            break;
        case 'O': /* ALTER DOMAIN SET NOT NULL */
            address =
                AlterDomainNotNull(stmt->typeName,
                                   true);

            break;
        case 'C': /* ADD CONSTRAINT */
            address =
                AlterDomainAddConstraint(stmt->typeName,
                                         stmt->def,
                                         &secondaryObject);

            break;
        case 'X': /* DROP CONSTRAINT */
            address =
                AlterDomainDropConstraint(stmt->typeName,
                                         stmt->name,
                                         stmt->behavior,
                                         stmt->missing_ok);

            break;
        case 'V': /* VALIDATE CONSTRAINT */
            address =

```



```

        AlterDomainValidateConstraint(stmt->typeName,
                                     stmt->name);

        break;
    default: /* oops */
        elog(ERROR, "unrecognized alter domain type: %d",
              (int) stmt->subtype);
        break;
    }
}
break;

/*
 * ***** object creation / destruction *****
 */
case T_DefineStmt:
{
    DefineStmt *stmt = (DefineStmt *) parsetree;

    switch (stmt->kind)
    {
        case OBJECT_AGGREGATE:
            address =
                DefineAggregate(pstate, stmt->defnames, stmt->args,
                               stmt->oldstyle,
                               stmt->definition,
                               stmt->replace);

            break;
        case OBJECT_OPERATOR:
            Assert(stmt->args == NIL);
            address = DefineOperator(stmt->defnames,
                                    stmt->definition);

            break;
        case OBJECT_TYPE:
            Assert(stmt->args == NIL);
            address = DefineType(pstate,
                                stmt->defnames,
                                stmt->definition);

            break;
        case OBJECT_TSPARSER:
            Assert(stmt->args == NIL);
            address = DefineTSParser(stmt->defnames,
                                    stmt->definition);

            break;
        case OBJECT_TSDICTIONARY:
            Assert(stmt->args == NIL);
            address = DefineTSDictionary(stmt->defnames,
                                        stmt->definition);

            break;
        case OBJECT_TSTEMPLATE:

```

```

        Assert(stmt->args == NIL);
        address = DefineTSTemplate(stmt->defnames,
                                   stmt->definition);

        break;
    case OBJECT_TSCONFIGURATION:
        Assert(stmt->args == NIL);
        address = DefineTSConfiguration(stmt->defnames,
                                       stmt->definition,
                                       &secondaryObject);

        break;
    case OBJECT_COLLATION:
        Assert(stmt->args == NIL);
        address = DefineCollation(pstate,
                                  stmt->defnames,
                                  stmt->definition,
                                  stmt->if_not_exists);

        break;
    default:
        elog(ERROR, "unrecognized define stmt type: %d",
             (int) stmt->kind);
        break;
    }
}
break;

case T_IndexStmt: /* CREATE INDEX */
{
    IndexStmt *stmt = (IndexStmt *) parsetree;
    Oid        relid;
    LOCKMODE   lockmode;

    if (stmt->concurrent)
        PreventInTransactionBlock(isTopLevel,
                                   "CREATE INDEX CONCURRENTLY");

    /*
     * Look up the relation OID just once, right here at the
     * beginning, so that we don't end up repeating the name
     * lookup later and latching onto a different relation
     * partway through. To avoid lock upgrade hazards, it's
     * important that we take the strongest lock that will
     * eventually be needed here, so the lockmode calculation
     * needs to match what DefineIndex() does.
     */
    lockmode = stmt->concurrent ? ShareUpdateExclusiveLock
                               : ShareLock;
    relid =
        RangeVarGetRelidExtended(stmt->relation, lockmode,
                                0,

```

```

        RangeVarCallbackOwnsRelation,
        NULL);

/*
 * CREATE INDEX on partitioned tables (but not regular
 * inherited tables) recurses to partitions, so we must
 * acquire locks early to avoid deadlocks.
 *
 * We also take the opportunity to verify that all
 * partitions are something we can put an index on, to
 * avoid building some indexes only to fail later.
 */
if (stmt->relation->inh &&
    get_rel_relkind(relid) == RELKIND_PARTITIONED_TABLE)
{
    ListCell    *lc;
    List        *inheritors = NIL;

    inheritors = find_all_inheritors(relid, lockmode, NULL);
    foreach(lc, inheritors)
    {
        char        relkind = get_rel_relkind(lfirst_oid(lc));

        if (relkind != RELKIND_RELATION &&
            relkind != RELKIND_MATVIEW &&
            relkind != RELKIND_PARTITIONED_TABLE &&
            relkind != RELKIND_FOREIGN_TABLE)
            elog(ERROR, "unexpected relkind \"%c\" on partition \"%s\"",
                 relkind, stmt->relation->relname);

        if (relkind == RELKIND_FOREIGN_TABLE &&
            (stmt->unique || stmt->primary))
            ereport(ERROR,
                    (errcode(ERRCODE_WRONG_OBJECT_TYPE),
                     errmsg("cannot create unique index on partitioned table
%s",
                             stmt->relation->relname),
                     errdetail("Table \"%s\" contains partitions that are
foreign tables.",
                             stmt->relation->relname))));
    }
    list_free(inheritors);
}

/* Run parse analysis ... */
stmt = transformIndexStmt(relid, stmt, queryString);

/* ... and do it */
EventTriggerAlterTableStart(parsetree);

```

```

address =
    DefineIndex(relid, /* OID of heap relation */
                stmt,
                InvalidOid, /* no predefined OID */
                InvalidOid, /* no parent index */
                InvalidOid, /* no parent constraint */
                false, /* is_alter_table */
                true, /* check_rights */
                true, /* check_not_in_use */
                false, /* skip_build */
                false); /* quiet */

/*
 * Add the CREATE INDEX node itself to stash right away;
 * if there were any commands stashed in the ALTER TABLE
 * code, we need them to appear after this one.
 */
EventTriggerCollectSimpleCommand(address, secondaryObject,
                                  parsetree);
commandCollected = true;
EventTriggerAlterTableEnd();
}
break;

case T_CreateExtensionStmt:
    address = CreateExtension(pstate, (CreateExtensionStmt *) parsetree);
    break;

case T_AlterExtensionStmt:
    address = ExecAlterExtensionStmt(pstate, (AlterExtensionStmt *)
parsetree);
    break;

case T_AlterExtensionContentsStmt:
    address = ExecAlterExtensionContentsStmt((AlterExtensionContentsStmt *)
parsetree,
                                              &secondaryObject);
    break;

case T_CreateFdwStmt:
    address = CreateForeignDataWrapper((CreateFdwStmt *) parsetree);
    break;

case T_AlterFdwStmt:
    address = AlterForeignDataWrapper((AlterFdwStmt *) parsetree);
    break;

case T_CreateForeignServerStmt:
    address = CreateForeignServer((CreateForeignServerStmt *) parsetree);

```

```

        break;

    case T_AlterForeignServerStmt:
        address = AlterForeignServer((AlterForeignServerStmt *) parsetree);
        break;

    case T_CreateUserMappingStmt:
        address = CreateUserMapping((CreateUserMappingStmt *) parsetree);
        break;

    case T_AlterUserMappingStmt:
        address = AlterUserMapping((AlterUserMappingStmt *) parsetree);
        break;

    case T_DropUserMappingStmt:
        RemoveUserMapping((DropUserMappingStmt *) parsetree);
        /* no commands stashed for DROP */
        commandCollected = true;
        break;

    case T_ImportForeignSchemaStmt:
        ImportForeignSchema((ImportForeignSchemaStmt *) parsetree);
        /* commands are stashed inside ImportForeignSchema */
        commandCollected = true;
        break;

    case T_CompositeTypeStmt: /* CREATE TYPE (composite) */
    {
        CompositeTypeStmt *stmt = (CompositeTypeStmt *) parsetree;

        address = DefineCompositeType(stmt->typevar,
                                      stmt->coldeflist);
    }
    break;

    case T_CreateEnumStmt: /* CREATE TYPE AS ENUM */
        address = DefineEnum((CreateEnumStmt *) parsetree);
        break;

    case T_CreateRangeStmt: /* CREATE TYPE AS RANGE */
        address = DefineRange((CreateRangeStmt *) parsetree);
        break;

    case T_AlterEnumStmt: /* ALTER TYPE (enum) */
        address = AlterEnum((AlterEnumStmt *) parsetree);
        break;

    case T_ViewStmt: /* CREATE VIEW */
        EventTriggerAlterTableStart(parsetree);

```

```

address = DefineView((ViewStmt *) parsetree, queryString,
                    pstmt->stmt_location, pstmt->stmt_len);
EventTriggerCollectSimpleCommand(address, secondaryObject,
                                parsetree);
/* stashed internally */
commandCollected = true;
EventTriggerAlterTableEnd();
break;

case T_CreateFunctionStmt: /* CREATE FUNCTION */
    address = CreateFunction(pstate, (CreateFunctionStmt *) parsetree);
    break;

case T_AlterFunctionStmt: /* ALTER FUNCTION */
    address = AlterFunction(pstate, (AlterFunctionStmt *) parsetree);
    break;

case T_RuleStmt: /* CREATE RULE */
    address = DefineRule((RuleStmt *) parsetree, queryString);
    break;

case T_CreateSeqStmt:
    address = DefineSequence(pstate, (CreateSeqStmt *) parsetree);
    break;

case T_AlterSeqStmt:
    address = AlterSequence(pstate, (AlterSeqStmt *) parsetree);
    break;

case T_CreateTableAsStmt:
    address = ExecCreateTableAs((CreateTableAsStmt *) parsetree,
                                queryString, params, queryEnv,
                                completionTag);
    break;

case T_RefreshMatViewStmt:

    /*
     * REFRESH CONCURRENTLY executes some DDL commands internally.
     * Inhibit DDL command collection here to avoid those commands
     * from showing up in the deparsed command queue. The refresh
     * command itself is queued, which is enough.
     */
    EventTriggerInhibitCommandCollection();
    PG_TRY();
    {
        address = ExecRefreshMatView((RefreshMatViewStmt *) parsetree,
                                    queryString, params, completionTag);
    }
}

```

```

PG_CATCH();
{
    EventTriggerUndoInhibitCommandCollection();
    PG_RE_THROW();
}
PG_END_TRY();
EventTriggerUndoInhibitCommandCollection();
break;

case T_CreateTrigStmt:
    address = CreateTrigger((CreateTrigStmt *) parsetree,
                            queryString, InvalidOid, InvalidOid,
                            InvalidOid, InvalidOid, InvalidOid,
                            InvalidOid, NULL, false, false);
    break;

case T_CreatePLangStmt:
    address = CreateProceduralLanguage((CreatePLangStmt *) parsetree);
    break;

case T_CreateDomainStmt:
    address = DefineDomain((CreateDomainStmt *) parsetree);
    break;

case T_CreateConversionStmt:
    address = CreateConversionCommand((CreateConversionStmt *) parsetree);
    break;

case T_CreateCastStmt:
    address = CreateCast((CreateCastStmt *) parsetree);
    break;

case T_CreateOpClassStmt:
    DefineOpClass((CreateOpClassStmt *) parsetree);
    /* command is stashed in DefineOpClass */
    commandCollected = true;
    break;

case T_CreateOpFamilyStmt:
    address = DefineOpFamily((CreateOpFamilyStmt *) parsetree);
    break;

case T_CreateTransformStmt:
    address = CreateTransform((CreateTransformStmt *) parsetree);
    break;

case T_AlterOpFamilyStmt:
    AlterOpFamily((AlterOpFamilyStmt *) parsetree);
    /* commands are stashed in AlterOpFamily */

```

```

    commandCollected = true;
    break;

case T_AlterTSDictionaryStmt:
    address = AlterTSDictionary((AlterTSDictionaryStmt *) parsetree);
    break;

case T_AlterTSConfigurationStmt:
    AlterTSConfiguration((AlterTSConfigurationStmt *) parsetree);

    /*
     * Commands are stashed in MakeConfigurationMapping and
     * DropConfigurationMapping, which are called from
     * AlterTSConfiguration
     */
    commandCollected = true;
    break;

case T_AlterTableMoveAllStmt:
    AlterTableMoveAll((AlterTableMoveAllStmt *) parsetree);
    /* commands are stashed in AlterTableMoveAll */
    commandCollected = true;
    break;

case T_DropStmt:
    ExecDropStmt((DropStmt *) parsetree, isTopLevel);
    /* no commands stashed for DROP */
    commandCollected = true;
    break;

case T_RenameStmt:
    address = ExecRenameStmt((RenameStmt *) parsetree);
    break;

case T_AlterObjectDependsStmt:
    address =
        ExecAlterObjectDependsStmt((AlterObjectDependsStmt *) parsetree,
                                    &secondaryObject);
    break;

case T_AlterObjectSchemaStmt:
    address =
        ExecAlterObjectSchemaStmt((AlterObjectSchemaStmt *) parsetree,
                                   &secondaryObject);
    break;

case T_AlterOwnerStmt:
    address = ExecAlterOwnerStmt((AlterOwnerStmt *) parsetree);
    break;

```



```

case T_AlterOperatorStmt:
    address = AlterOperator((AlterOperatorStmt *) parsetree);
    break;

case T_CommentStmt:
    address = CommentObject((CommentStmt *) parsetree);
    break;

case T_GrantStmt:
    ExecuteGrantStmt((GrantStmt *) parsetree);
    /* commands are stashed in ExecGrantStmt_oids */
    commandCollected = true;
    break;

case T_DropOwnedStmt:
    DropOwnedObjects((DropOwnedStmt *) parsetree);
    /* no commands stashed for DROP */
    commandCollected = true;
    break;

case T_AlterDefaultPrivilegesStmt:
    ExecAlterDefaultPrivilegesStmt(pstate, (AlterDefaultPrivilegesStmt *)
parsetree);
    EventTriggerCollectAlterDefPrivs((AlterDefaultPrivilegesStmt *)
parsetree);
    commandCollected = true;
    break;

case T_CreatePolicyStmt: /* CREATE POLICY */
    address = CreatePolicy((CreatePolicyStmt *) parsetree);
    break;

case T_AlterPolicyStmt: /* ALTER POLICY */
    address = AlterPolicy((AlterPolicyStmt *) parsetree);
    break;

case T_SecLabelStmt:
    address = ExecSecLabelStmt((SecLabelStmt *) parsetree);
    break;

case T_CreateAmStmt:
    address = CreateAccessMethod((CreateAmStmt *) parsetree);
    break;

case T_CreatePublicationStmt:
    address = CreatePublication((CreatePublicationStmt *) parsetree);
    break;

```

[illegible]

```
        EventTriggerSQLDrop(parsetree);
        EventTriggerDDLCommandEnd(parsetree);
    }
}
PG_CATCH();
{
    if (needCleanup)
        EventTriggerEndCompleteQuery();
    PG_RE_THROW();
}
PG_END_TRY();

if (needCleanup)
    EventTriggerEndCompleteQuery();
}
```